

Experiment 12

Aim: To create a Lambda function which will log "An Image has been added" once you add an object to a specific bucket in S3.

Theory:

Creating a system to log activities when an image is added to an S3 bucket involves integrating Amazon S3 with AWS Lambda.

Amazon S3:

Amazon S3 (Simple Storage Service) is a service offered by AWS that provides object storage through a web service interface. It's designed to store and retrieve any amount of data from anywhere. You can use S3 to store images, videos, backups, data logs, and more.

AWS Lambda:

AWS Lambda is a serverless compute service that allows you to run code in response to events without provisioning or managing servers. You write your code and set up a trigger, and Lambda takes care of the rest. This means that when a specified event occurs, such as an object being added to an S3 bucket, the Lambda function is automatically invoked.

By setting up a Lambda function to trigger on new uploads to a specific S3 bucket, you can automate logging activities. This function will capture the event, process it, and log the message "An Image has been added." It ensures that every new upload is tracked efficiently and that you have a record of these actions.

This kind of setup is highly scalable, reliable, and costeffective, leveraging AWS's robust infrastructure. It's particularly useful for applications that require automated monitoring and logging of uploads for auditing or notification purposes.

Implementation:

1. Create an S3 Bucket: First, create an S3 bucket that will store the objects. This bucket will act as the trigger source for the Lambda function.

aws

Services

Search

[Alt+S]

N. V

Amazon S3 > Buckets > Create bucket

Create bucket [Info](#)

Buckets are containers for data stored in S3.

General configuration

AWS Region

US East (N. Virginia) us-east-1

Bucket type

☒ General purpose

Recommended for most use cases and access patterns. General purpose buckets are the original S3 bucket type. They allow a mix of storage classes that redundantly store objects across multiple Availability Zones.

☐ Directory

Recommended for low-latency use cases. These buckets use only the S3 Express One Zone storage class, which provides faster processing of data within a single Availability Zone.

Bucket name

myawsbucket

Bucket name must be unique within the global namespace and follow the bucket naming rules. [See rules for bucket naming](#)

Copy settings from existing bucket - optional

Only the bucket settings in the following configuration are copied.

aws

Services

Search

[Alt+S]

N. Virginia

General configuration

AWS Region

US East (N. Virginia) us-east-1

Bucket type

☒ General purpose

Recommended for most use cases and access patterns. General purpose buckets are the original S3 bucket type. They allow a mix of storage classes that redundantly store objects across multiple Availability Zones.

☐ Directory

Recommended for low-latency use cases. These buckets use only the S3 Express One Zone storage class, which provides faster processing of data within a single Availability Zone.

Bucket name

ronakbucket123

Bucket name must be unique within the global namespace and follow the bucket naming rules. [See rules for bucket naming](#)

Copy settings from existing bucket - optional

Only the bucket settings in the following configuration are copied.

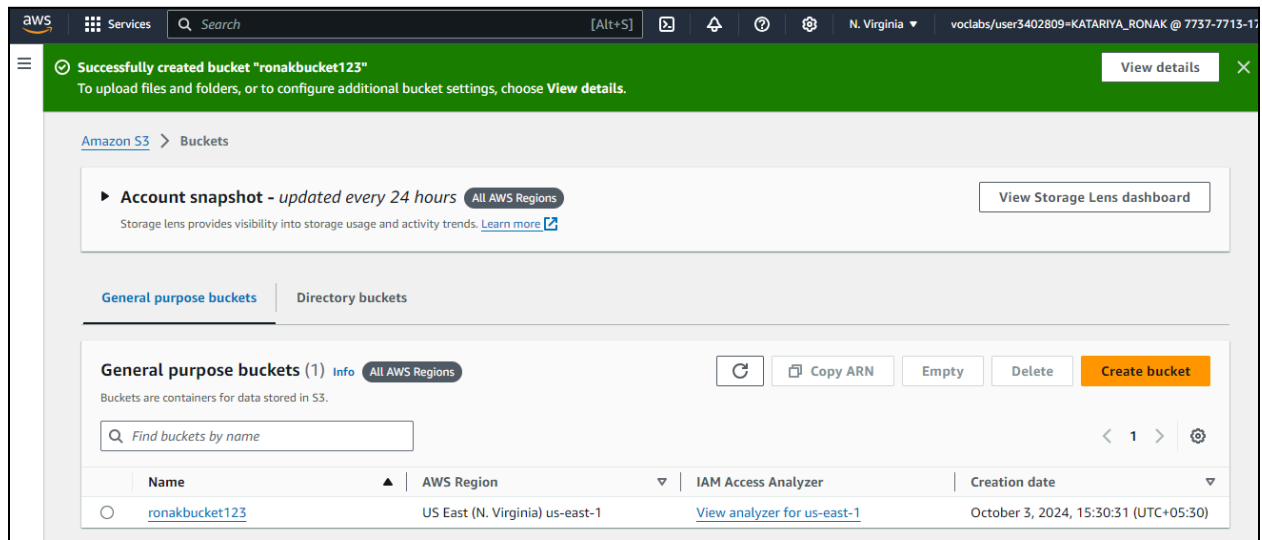
Choose bucket

Format: s3://bucket/prefix

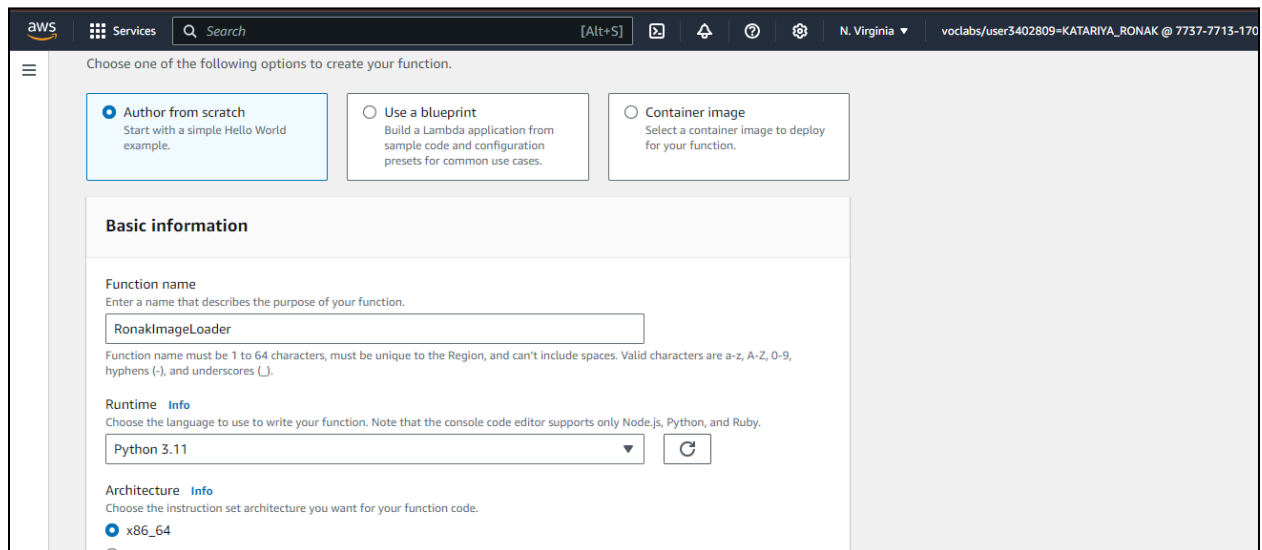
Object Ownership

[Info](#)

Bucket created successfully:



2. Create the Lambda Function: Set up a new Lambda function using AWS Lambda's console. You can choose a runtime environment like Python, Node.js, or Java.



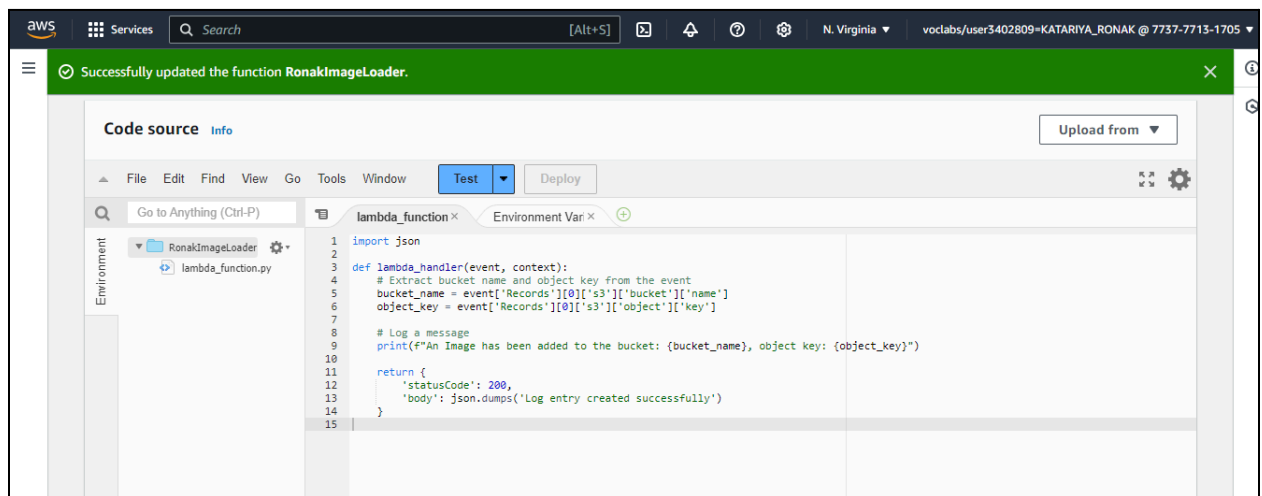
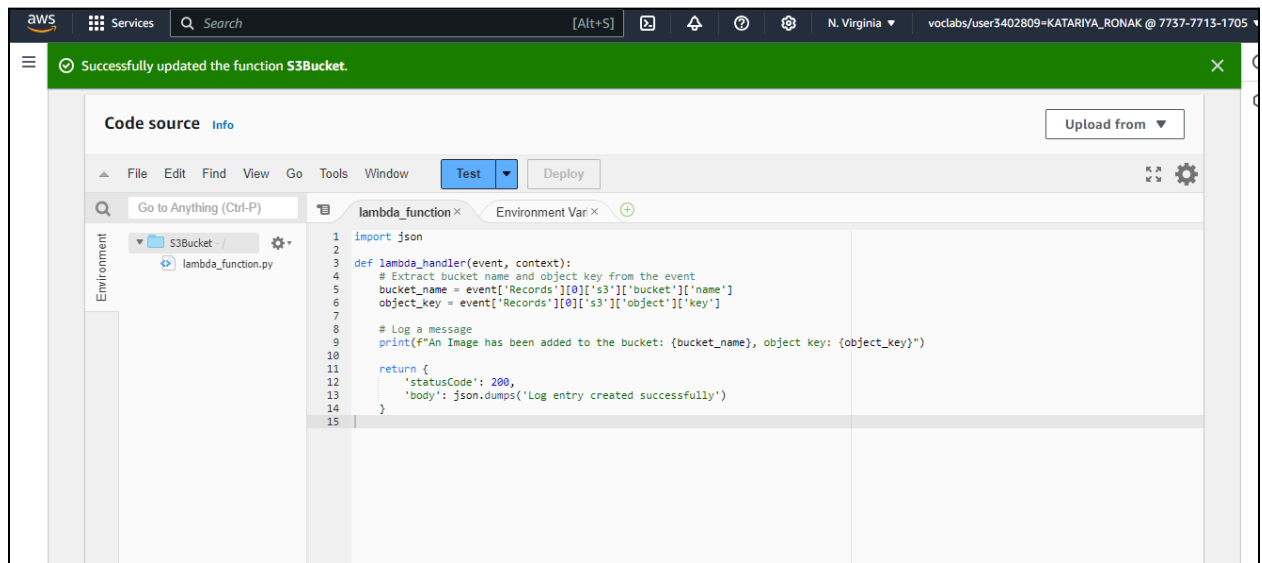
Write code that logs a message like "An Image has been added" when triggered
This is the code we have to add:

```
import json
import
logging
# Set up logging
logger = logging.getLogger()
logger.setLevel(logging.INFO)
def lambda_handler(event, context):
```

```

# Extract bucket name and object key from the S3
event for record in event['Records']:
    bucket = record['s3']['bucket']['name']
    key = record['s3']['object']['key']
# Log a message
logger.info(f"An image has been added to bucket {bucket}, object key: {key}")
# Check if the uploaded file is an image (you can adjust the file types
here) if key.lower().endswith(('.png', '.jpg', '.jpeg', '.gif')):
    logger.info("Image Uploaded
successfully") else:
    logger.info("A non-image file has been added")
return {
'statusCode': 200,
'body': json.dumps('Event processed')
}

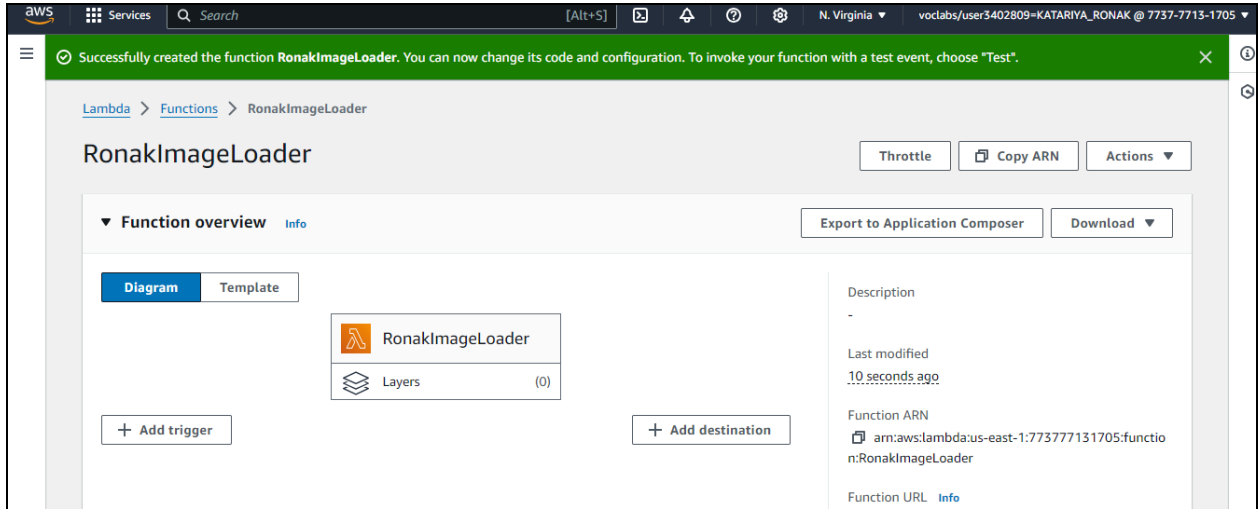
```



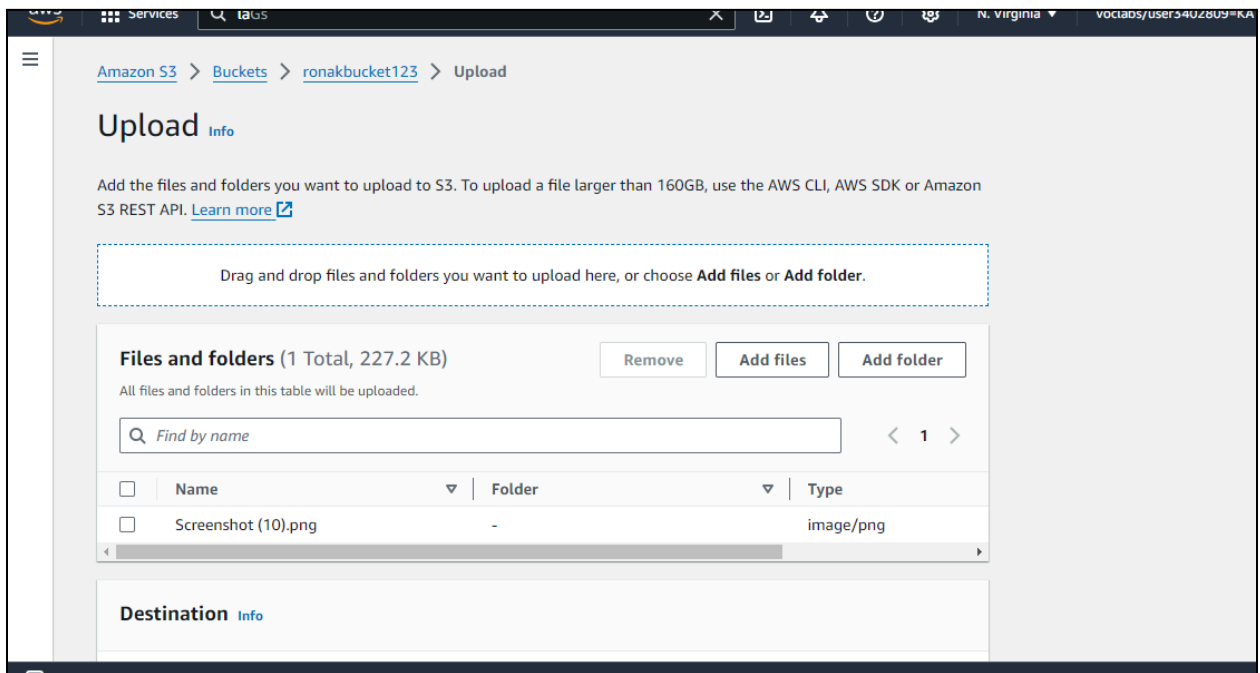
Configure S3 Trigger: Link the S3 bucket to the Lambda function by setting up a trigger. Specify that the function should be triggered when an object is created in the bucket (e.g., when an image is uploaded).

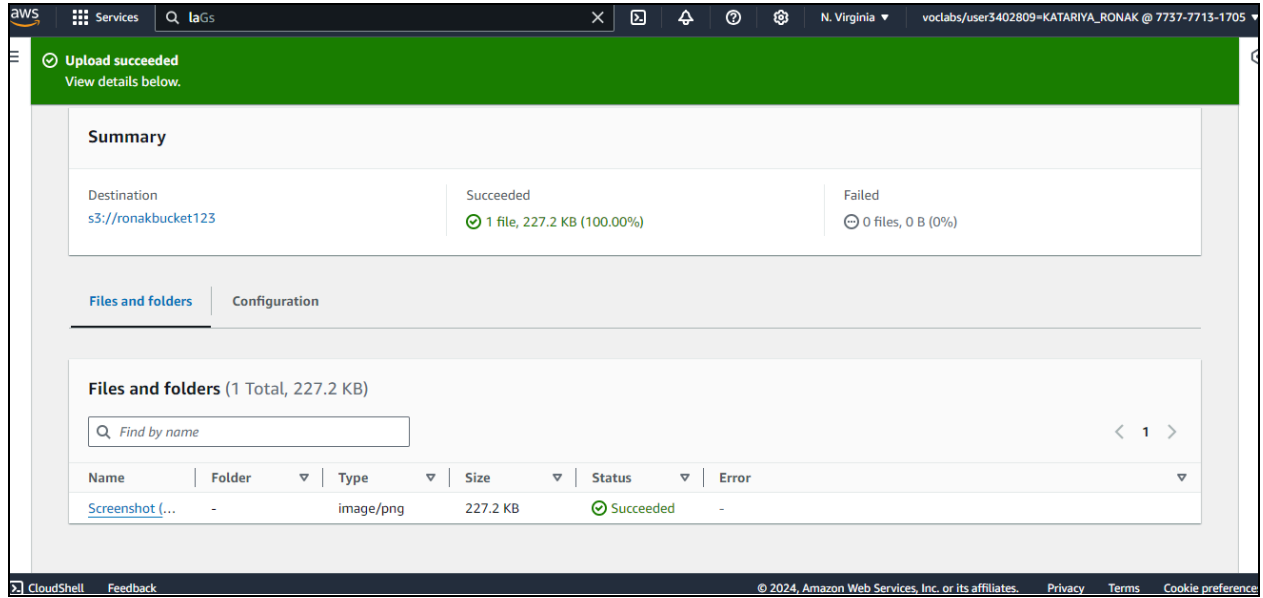
The screenshot shows the 'Add trigger' configuration page in the AWS Lambda console. The 'Trigger configuration' section is active, showing a dropdown menu for 'S3' with the text 'aws asynchronous storage'. Below this, the 'Bucket' section prompts the user to 'Choose or enter the ARN of an S3 bucket that serves as the event source. The bucket must be in the same region as the function.' A search bar contains 's3/ronakbucket123', and the 'Bucket region' is set to 'us-east-1'. The 'Event types' section instructs the user to 'Select the events that you want to have trigger the Lambda function. You can optionally set up a prefix or suffix for an event. However, for each bucket, individual events cannot have multiple configurations with overlapping prefixes or suffixes that could match the same object key.' A dropdown menu is shown with 'All object create events' selected.

The screenshot shows the 'S3Bucket' function overview page in the AWS Lambda console. A green notification bar at the top states: 'The trigger ronakbucket123 was successfully added to function S3Bucket. The function is now receiving events from the trigger.' The 'Function overview' section includes a 'Diagram' tab and a 'Template' tab. The diagram shows a box for 'S3Bucket' with a 'Layers' section indicating '(0)' layers. Below the diagram is a box for 'S3' with a '+ Add trigger' button. To the right of the diagram is a '+ Add destination' button. On the right side of the page, there is a 'Description' section with a '-' icon, a 'Last modified' section showing '7 days ago', a 'Function ARN' section with the value 'arn:aws:lambda:us-east-1:773777131705:function:S3Bucket', and a 'Function URL' section with an 'Info' link.

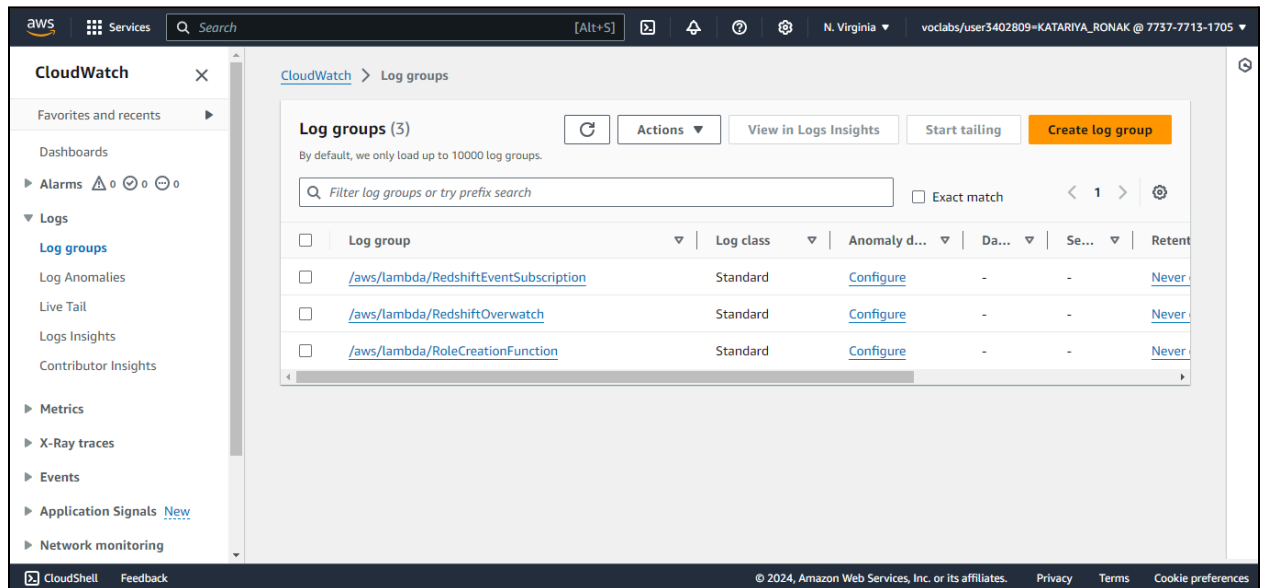


Upload an object (e.g., an image) to the S3 bucket to test the trigger





Test the Setup: Upload an object (e.g., an image) to the S3 bucket to test the trigger. The Lambda function should execute and log the message "An Image has been added" in AWS CloudWatch Logs



Log events

You can use the filter bar below to search for and match terms, phrases, or values in your log events. [Learn more about filter patterns](#)

Actions

Start tailing

Create metric filter

Clear

1m

30m

1h

12h

Custom

Local

Display

Timestamp	Message
No older events at this moment. Retry	
2023-10-07T17:49:20.002+05:30	INIT_START Runtime Version: python:3.11.v14 Runtime Version ARN: arn:aws:lambda:ap-south-1::runtime:9c87c21a94b293e1a306aad2c23c...
2023-10-07T17:49:20.110+05:30	START RequestId: a189471e-867f-4db4-824f-2b602f956879 Version: \$LATEST
2023-10-07T17:49:20.111+05:30	END RequestId: a189471e-867f-4db4-824f-2b602f956879
2023-10-07T17:49:20.111+05:30	REPORT RequestId: a189471e-867f-4db4-824f-2b602f956879 Duration: 1.25 ms Billed Duration: 2 ms Memory Size: 128 MB Max Memory Us...
No newer events at this moment. <i>Auto retry paused.</i> Resume	

Conclusion:

Creating a Lambda function to log “An Image has been added” upon S3 uploads showcases AWS's serverless capabilities. This integration automates event handling, enhances operational efficiency, and reduces manual oversight, allowing developers to focus on application logic while ensuring responsive and scalable cloud applications.